Functions - part 2
STAT 133

Gaston Sanchez

Department of Statistics, UC–Berkeley

gastonsanchez.com
github.com/gastonstat/stat133
Course web: gastonsanchez.com/teaching/stat133

Functions

R comes with many functions and packages that let us perform a wide variety of tasks. Sometimes, however, there's no function to do what we want to achieve. In these cases we need to create our own functions.

Writing Functions

Writing Functions

Writing Functions

- ▶ Choose meaningful names of functions
- ▶ Preferably a verb
- ▶ Choose meaningful names of arguments
- ▶ Think about the users (who will use the function)
- ▶ Think about extreme cases
- ▶ If a function is too long, maybe you need to split it

Names of functions

Avoid this:

```r
f <- function(x, y) {
  x + y
}
```

This is better

```r
add <- function(x, y) {
  x + y
}
```

Names of arguments

Give meaningful names to arguments:

```
# Avoid this
area_rect <- function(x, y) {
  x * y
}
```

This is better

```
area_rect <- function(length, width) {
  length * width
}
```

Meaningful Names to Arguments

Avoid this:

```r
# what does this function do?
ci <- function(p, r, n, ti) {
  p * (1 + r/p)^(ti * p)
}
```

Meaningful Names to Arguments

Avoid this:

```
# what does this function do?
ci <- function(p, r, n, ti) {
  p * (1 + r/p)^(ti * p)
}
```

This is better:

```
# OK
compound_interest <- function(principal, rate, periods, time) {
  principal * (1 + rate/periods)^(time * periods)
}
```

Meaningful Names to Arguments

```r
# names of arguments
compound_interest <- function(principal = 1, rate = 0.01,
                              periods = 1, time = 1)
{
  principal * (1 + rate/periods)^(time * periods)
}

compound_interest(principal = 100, rate = 0.05,
                  periods = 5, time = 1)

compound_interest(rate = 0.05, periods = 5,
                  time = 1, principal = 100)

compound_interest(rate = 0.05, time = 1,
                  periods = 5, principal = 100)
```

Describing functions

Also add a short description of what the arguments should be like. In this case, the description is outside the function

```r
# function for adding two numbers
# x: number
# y: number
add <- function(x, y) {
  x + y
}
```

Describing functions

In this case, the description is inside the function

```r
add <- function(x, y) {
  # function for adding two numbers
  # x: number
  # y: number
  x + y
}
```

Describing functions

```r
# description of arguments
compound_interest <- function(principal = 1, rate = 0.01,
                              periods = 1, time = 1)
{
  # principal = Principal Amount
  # rate = Annual Nominal Interest Rate as a decimal
  # time = Time Involved in years
  # periods = number of compounding periods per unit time
  principal * (1 + rate/periods)^(time * periods)
}
```

Binary Operators

Binary Operators

- One type of functions very common in R are **binary operators**, eg:
    - `2 + 5` (sum)
    - `3 / 2` (division)
    - `a %in% b` (value matching)
    - `X %*% Y` (matrix multiplication)
- Binary operators are actually functions
- These functions take two inputs—hence the term *binary*
- It is possible to define your own binary operators

Binary Operators

Example:

```
# addition operator
2 + 3

# equivalent to
'+'(2, 3)
```

Binary Operators

```r
# binary operator
"%p%" <- function(x, y) {
  paste(x, y, sep = " ")
}

'good' %p% 'morning'

## [1] "good morning"
```

How to create a binary operator?

- A binary operator is defined as one or more characters surrounded by percent symbols %

- When defining the function, the entire name must be quoted

- Include two arguments

- As usual, avoid using names of existing operators:
  - "%%", %*%, %/%, %o%, %in%

Another example

Here's another example:

```r
# binary operator
"%u%" <- function(x, y) {
  union(x, y)
}

1:5 %u% c(1, 3, 5, 7, 9)

## [1] 1 2 3 4 5 7 9
```

Lazy Evaluation

Lazy Evaluation

Arguments to functions are evaluated lazily, that is, they are evaluated only as needed:

```
g <- function(a, b) {
  a * a * a
}

g(2)

## [1] 8
```

g() never uses the argument b, so calling g(2) does not produce an error

Lazy Evaluation

Another example

```r
g <- function(a, b) {
  print(a)
  print(b)
}

g(2)

## [1] 2

## Error in print(b):  argument "b" is missing, with no
default
```

Notice that 2 got printed before the error was triggered. This is
because b did not have to be evaluated until after print(a)

Messages

There are two main functions for generating warnings and errors:

- ► `stop()`
- ► `warning()`

There's also the `stopifnot()` function

Stop Execution

Use stop() to stop the execution of a function (this will raise an error)

```r
meansd <- function(x, na.rm = FALSE) {
  if (!is.numeric(x)) {
    stop("x is not numeric")
  }
  # output
  c(mean = mean(x, na.rm = na.rm),
    sd = sd(x, na.rm = na.rm))
}
```

Stop Execution

Use `stop()` to stop the execution of a function (this will raise an error)

```r
# ok
meansd(c(4, 5, 3, 1, 2))

##     mean       sd
## 3.000000 1.581139

# this causes an error
meansd(c('a', 'b', 'c'))

## Error in meansd(c("a", "b", "c")):  x is not numeric
```

Warning Messages

Use warning() to show a warning message

```r
meansd <- function(x, na.rm = FALSE) {
  if (!is.numeric(x)) {
    warning("non-numeric input coerced to numeric")
    x <- as.numeric(x)
  }
  # output
  c(mean = mean(x, na.rm = na.rm),
    sd = sd(x, na.rm = na.rm))
}
```

A warning is useful when we don't want to stop the execution, but we still want to show potential problems

Warning Messages

Use warning() to show a warning message

```
# ok
meansd(c(4, 5, 3, 1, 2))

##     mean       sd
## 3.000000 1.581139

# this causes a warning
meansd(c(TRUE, FALSE, TRUE, FALSE))

## Warning in meansd(c(TRUE, FALSE, TRUE, FALSE)):
## non-numeric input coerced to numeric

##      mean        sd
## 0.5000000 0.5773503
```

Stop Execution

`stopifnot()` ensures the truth of expressions:

```r
meansd <- function(x, na.rm = FALSE) {
  stopifnot(is.numeric(x))
  # output
  c(mean = mean(x, na.rm = na.rm),
    sd = sd(x, na.rm = na.rm))
}

meansd('hello')

## Error:  is.numeric(x) is not TRUE
```

Environments and Functions

Consider this example

```
w <- 10

f <- function(y) {
  d <- 5
  h <- function() {
    d * (w + y)
  }
  return(h())
}

f(2)

## [1] 60
```

How / Why does f() work?

Consider this other example

```
w <- 10

f <- function(y) {
  d <- 5
  return(h())
}

f(2)

## Error in f(2):  could not find function "h"
```

Why f() does not work?

Environments

- All the variables that we create need to be stored somewhere

- The place where they are stored is called an **environment**

- R works with enviroments, all of which are in (virtual) memory

- Usually, we don't need to explicitly deal with environments

- Environments are nested

Global Environment

- The user workspace is the **global environment**
- The global environment is the **top level** environment
- It is formally referred to as R_GlobalEnv
- Variables defined in the global environment can be seen from anywhere
- The contents of the global environment are listed with `ls()`

```r
# top level environment
environment()

## <environment: R_GlobalEnv>
```

Searching objects

- When R tries to bind a value to a symbol, it searches through a series of environments to find the appropriate value
- To retrieve the value of an object the order is:
- Search the current environment
- Search the global environment for a symbol name matching the one requested
- Search the namespaces of each of the packages on the search list: `search()`

Environments and Functions

- A function consists not only of its arguments and body but also of its *environment*

- An environment is made up of the collection of objects present at the time the function comes into existence

- When a function is created by evaluating the corresponding expression, the current environment is recorded as a property of the function

Let's go back to our first example

```r
w <- 10

f <- function(y) {
  d <- 5
  h <- function() {
    d * (w + y)
  }
  return(h())
}

f(2)

## [1] 60
```

How does f() work?

Let's see the environments

```r
w <- 10    # variable (in global environment)

# a function (in global environment)
f <- function(y) {
  d <- 5      # local variable
  h <- function() {    # subfunction
    d * (w + y)     # w is a free variable
  }
  return(h())
}

environment(f)

## <environment: R_GlobalEnv>
```

Function Environment

- `w` is a global variable (in global environment)
- `f()` is a function in the global environment
- `d` is a local variable—local to `f()`
- `h()` is a subfunction—local to `f()`
- `w` is not an argument but a free variable

Let's see the environments

```r
f <- function(y) {
  d <- 5
  h <- function() {
    d * (w + y)
  }
  print(environment(h))   # h()'s environment
  return(h())
}

environment(f)

## <environment: R_GlobalEnv>

f(2)

## <environment: 0x7fdf54976778>
## [1] 60
```

Your turn

```r
# When executed, what does g() return?
x <- 5

g <- function(x = FALSE) {
    y <- 10
    list(x = x, y = y)
}
g()
```

A) x = 5, y = 10
B) x = 0, y = 10
C) x = FALSE, y = 10
D) x = FALSE, y = 5

Variable's Scope

- A variable's **scope** is the set of places from which you can see the variable
- R will try to find a variable in the current environment
- If it doesn't find them it will look in the parent environment
- And then that environment's parent
- And so on until it reaches the global environment

Variable's Scope

- A variable's **scope** is the set of places from which you can see the variable
- R will try to find a variable in the current environment
- If it doesn't find them it will look in the parent environment
- And then that environment's parent
- And so on until it reaches the global environment

Variable Scope

- When we define a variable inside a function, the rest of the statements in that function will have access to that variable

Variable Scope

```r
f <- function(x) {
  y <- 1
  g <- function(x) {
    (x + y) / 2
  }
  g(x)
}

f(5)

## [1] 3
```

g() is a subfunction that have access to y in f's environment.

Variable Scope

```
f <- function(x) {
  y <- 1
  g(x)
}

g <- function(x) {
  (x + y) / 2
}

f(5)

## Error in g(x):  object 'y' not found
```

g() is a function that doesn't have access to y; g() can only
see things in the global environment

One more thing ...

Let's look at another exmaple

```
mean(1:5)

## [1] 3

mean

## function (x, ...)
## UseMethod("mean")
## <bytecode: 0x7fdf55ac2870>
## <environment: namespace:base>
```

One more thing ...

You can do things like this

```
# confusing but it works
mean <- 1:5
mean(mean)

## [1] 3
```

Some issues

You can also do things like this

```r
# not a good idea but you can do it
mean <- function(x) 2*x + 5

mean(1:5)

## [1]  7  9 11 13 15
```

It seems we've lost the original mean() function

The :: Operator

:: operator to the rescue

```r
# my mean
mean(1:5)

## [1]  7  9 11 13 15

# base mean
base::mean(1:5)

## [1] 3
```

Here we use the name espace base of the R package "base" to access the original mean()

Recovering original `mean()`

To recover the original `mean()` you have to remove the artificial mean

```
# remove my mean
rm(mean)

# now try it again
mean(1:5)

## [1] 3
```

Your Turn

Exercise

R has a function summary() that when applied on a numeric vector provides something like this:

```
summary(1:10)

##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    1.00    3.25    5.50    5.50    7.75   10.00
```

Create a describe() function that takes a numeric vector and returns: minimum, maximum, mean, and standard deviation

Exercise

First attempt

```
describe <- function(x) {
  x_min <- min(x)
  x_max <- max(x)
  x_mean <- mean(x)
  x_sd <- sd(x)
  return(c(x_min, x_max, x_mean, x_sd))
}

describe(1:10)

## [1]  1.00000 10.00000  5.50000  3.02765
```

Exercise

Second attempt (adding names)

```r
describe <- function(x) {
  x_min <- min(x)
  x_max <- max(x)
  x_mean <- mean(x)
  x_sd <- sd(x)
  values <- c(x_min, x_max, x_mean, x_sd)
  names(values) <- c("min", "max", "mean", "sd")
  return(values)
}

describe(1:10)

##      min      max     mean       sd
##  1.00000 10.00000  5.50000  3.02765
```

Exercise

Third attempt (using a list as output)

```
describe <- function(x) {
  list(
    min = min(x),
    max = max(x),
    mean = mean(x),
    sd = sd(x)
  )
}
```

```
describe(1:10)

## $min
## [1] 1
##
## $max
## [1] 10
##
## $mean
## [1] 5.5
##
## $sd
## [1] 3.02765
```

Exercise

Probability Density of the Normal Distribution:

$$f(x|\mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

Write a function that takes a value x (with parameters μ and σ) which computes the probability density distribution of the normal distribution

Exercise

Normal Distribution:

```
normal_dist <- function(x, mu = 0, sigma = 1) {
  constant <- 1 / (sigma * sqrt(2*pi))
  constant * exp(-((x - mu)^2) / (2 * sigma^2))
}

normal_dist(2)

## [1] 0.05399097
```

Argument Matching

```
normal_dist <- function(x, mu = 0, sigma = 1) {
  constant <- 1 / (sigma * sqrt(2*pi))
  constant * exp(-((x - mu)^2) / (2 * sigma^2))
}

normal_dist(2)
normal_dist(2, sigma = 3, mu = 1)
normal_dist(mu = 1, sigma = 3, 2)
normal_dist(mu = 1, 2, sigma = 3)
```

Argument Matching

R is "smart" enough in doing pattern matching with arguments' names (not recommended though)

```
normal_dist(2)

## [1] 0.05399097

normal_dist(2, m = 0, s = 1)

## [1] 0.05399097

normal_dist(2, sig = 1, m = 0)

## [1] 0.05399097
```