

Regular Expressions 1

STAT 133

Gaston Sanchez

Department of Statistics, UC–Berkeley

`gastonsanchez.com`

`github.com/gastonstat/stat133`

Course web: gastonsanchez.com/teaching/stat133

Motivation

Motivation

- ▶ So far we have seen some basic and intermediate functions for handling and working with text in R
- ▶ These functions are very useful functions and they allows us to do many interesting things.
- ▶ However, if we truly want to unleash the power of strings manipulation, we need to talk about *regular expressions*.

Motivation

##	name	gender	height	weight	born
## 1	Luke	male	1.72m	77gr	19BBY
## 2	Leia	Female	1.50m	49gr	19BBY
## 3	R2-D2	male	0.96m	32gr	33BBY
## 4	C-3P0	MALE	1.67m	75gr	112BBY

Motivation

##	name	gender	height	weight	born
## 1	Luke	male	1.72m	77gr	19BBY
## 2	Leia	Female	1.50m	49gr	19BBY
## 3	R2-D2	male	0.96m	32gr	33BBY
## 4	C-3P0	MALE	1.67m	75gr	112BBY

It is not uncommon to have datasets that need some cleaning and preprocessing

Motivation

Some common tasks

- ▶ finding pieces of text or characters that meet a certain pattern
- ▶ extracting pieces of text in non-standard formats
- ▶ transforming text into a uniform format
- ▶ resolving inconsistencies
- ▶ substituting certain characters
- ▶ splitting a piece of text into various parts

Regular Expressions

Regular Expressions

- ▶ A **regular expression** is a special text string for describing a certain amount of text.
- ▶ This “certain amount of text” receives the formal name of **pattern**.
- ▶ A regular expression is a *pattern that describes a set of strings*.
- ▶ It is common to abbreviate the term “regular expression” as **regex**

Regular Expressions

Simply put, working with regular expressions is nothing more than **pattern matching**

Regular Expressions

- ▶ Regex patterns consist of a combination of alphanumeric characters as well as special characters.
 - e.g. `[a-zA-Z0-9_]*`
- ▶ A regex pattern can be as simple as a single character
- ▶ But it can also be formed by several characters with a more complex structure

Basic Concepts

Regular expressions are constructed from 3 things:

- ▶ **Literal characters** are matched only by the character itself
- ▶ A **character class** is matched by any single member of the specified class
- ▶ **Modifiers** operate on literal characters, character classes, or combinations of the two.

Literal Characters

Consider the following text:

The quick brown fox jumps over the lazy dog

A basic regex can be something as simple as `fox`. The characters `fox` match the word “fox” in the sentence above.

Special Characters

Consider this other text:

One. Two. Three. And four* and Five!

Not all characters are matched literally. There are some characters that have a special meaning in regular expressions: `.` or `*` are some of these special characters.

About Regex

- ▶ “Regular Expressions” is not a full programming language
- ▶ Regex have a special syntax and instructions that you must learn
- ▶ Regular expressions are supported in a variety of forms on almost every computing platform
- ▶ R has functions and packages designed to work with regular expressions

Regex Tasks

Common operations

- ▶ **Identify** a match to a pattern
- ▶ **Locate** a pattern match (positions)
- ▶ **Replace** a matched pattern
- ▶ **Extract** a matched pattern

Identify a Match

Identify Matches

Functions for identifying match to a pattern:

- ▶ `grep(..., value = FALSE)`
- ▶ `grepl()`
- ▶ `str_detect()` ("stringr")

grep(..., value = FALSE)

```
# some text
text <- c("one word", "a sentence", "three two one")

# pattern
pat <- "one"

# default usage
grep(pat, text)

## [1] 1 3
```

grep1()

- ▶ `grep1()` is very similar to `grep()`
- ▶ The difference resides in that the output are not numeric indices, but logical
- ▶ You can think of `grep1()` as *grep-logical*

grep1()

```
# some text
text <- c("one word", "a sentence", "three two one")

# pattern
pat <- "one"

# default usage
grep1(pat, text)

## [1]  TRUE FALSE  TRUE
```

str_detect() in "stringr"

```
# some text
text <- c("one word", "a sentence", "three two one")

# pattern
pat <- "one"

# default usage
str_detect(text, pat)

## [1]  TRUE FALSE  TRUE
```

Locate a Match

Locate Matches

Functions for locating match to a pattern:

- ▶ `regexpr()`
- ▶ `gregexpr()`
- ▶ `str_locate()` ("stringr")
- ▶ `str_locate_all()` ("stringr")

regexpr()

- ▶ `regexpr()` is used to find exactly where the pattern is found in a given string
- ▶ it tells us which elements of the text vector actually contain the regex pattern, and
- ▶ identifies the position of the substring that is matched by the regex pattern

regexpr()

```
# some text
text <- c("one word", "a sentence", "three two one")

# default usage
regexpr(pattern = "one", text)

## [1] 1 -1 11
## attr(,"match.length")
## [1] 3 -1 3
## attr(,"useBytes")
## [1] TRUE
```

gregexpr()

- ▶ `gregexpr()` does practically the same thing as `regexpr()`
- ▶ identifies where a pattern is within a string vector, by searching each element separately
- ▶ The only difference is that `gregexpr()` has an output in the form of a list

gregexpr()

```
# some text
text <- c("one word", "a sentence", "three two one")

# default usage
gregexpr(pattern = "one", text)

## [[1]]
## [1] 1
## attr(,"match.length")
## [1] 3
## attr(,"useBytes")
## [1] TRUE
##
## [[2]]
## [1] -1
## attr(,"match.length")
## [1] -1
## attr(,"useBytes")
## [1] TRUE
##
## [[3]]
## [1] 11
## attr(,"match.length")
## [1] 3
## attr(,"useBytes")
## [1] TRUE
```

`str_locate()` in "stringr"

- ▶ `str_locate()` locates the position of the first occurrence of a pattern in a string
- ▶ it tells us which elements of the text vector actually contain the regex pattern, and
- ▶ identifies the position of the substring that is matched by the regex pattern

str_locate() in "stringr"

```
# some text
text <- c("one word", "a sentence", "three two one")

# default usage
str_locate(text, pattern = "one")
```

##	start	end
## [1,]	1	3
## [2,]	NA	NA
## [3,]	11	13

`str_locate_all()` in "stringr"

- ▶ `str_locate_all()` locates the position of ALL the occurrences of a pattern in a string
- ▶ it tells us which elements of the text vector actually contain the regex pattern, and
- ▶ the output is in the form of a list with as many elements as the number of elements in the examined vector

str_locate_all() in "stringr"

```
# some text
text <- c("one word", "a sentence", "one three two one")

# default usage
str_locate_all(text, pattern = "one")

## [[1]]
##      start end
## [1,]     1   3
##
## [[2]]
##      start end
##
## [[3]]
##      start end
## [1,]     1   3
## [2,]    15  17
```

Extract a Match

Extract Matches

Functions for extracting positions of matched pattern:

- ▶ `grep(..., value = TRUE)`
- ▶ `str_extract() ("stringr")`
- ▶ `str_extract_all() ("stringr")`

Extraction with `grep(..., value = TRUE)`

- ▶ `grep(..., value = TRUE)` allows us to do basic extraction
- ▶ Actually, it will extract the entire element that matches the pattern
- ▶ Sometimes this is not exactly what we want to do (it's better to use functions of "stringr")

grep(..., value = TRUE)

```
# some text
text <- c("one word", "a sentence", "one three two one")

# extract first one
grep(pattern = "one", text, value = TRUE)

## [1] "one word"          "one three two one"
```

Pattern extraction with `str_extract()`

- ▶ `str_extract()` extracts the first occurrence of the matched pattern
- ▶ It will only extract the matched pattern
- ▶ If no pattern is matched, then a missing value is returned

str_extract() in "stringr"

```
# some text
text <- c("one word", "a sentence", "one three two one")

# extract first one
str_extract(text, pattern = "one")

## [1] "one" NA      "one"
```

Pattern extraction with `str_extract_all()`

- ▶ `str_extract_all()` extracts ALL the occurrences of the matched pattern
- ▶ It will only extract the matched pattern
- ▶ If no pattern is matched, then a character vector of length zero is returned
- ▶ the output is in list format

str_extract_all() in "stringr"

```
# some text
text <- c("one word", "a sentence", "one three two one")

# extract all 'one's
str_extract_all(text, pattern = "one")

## [[1]]
## [1] "one"
##
## [[2]]
## character(0)
##
## [[3]]
## [1] "one" "one"
```

Replace a Match

Replace Matches

Functions for replacing a matched pattern:

- ▶ `sub()`
- ▶ `gsub()`
- ▶ `str_replace()` ("stringr")
- ▶ `str_replace_all()` ("stringr")

Replace first occurrence with `sub()`

About `sub()`

- ▶ `sub()` replaces the **first** occurrence of a pattern in a given text
- ▶ if there is more than one occurrence of the pattern in each element of a string vector, only the first one will be replaced.

Replacing with sub()

```
# string
Rstring <- c("The R Foundation",
             "for Statistical Computing",
             "R is FREE software",
             "R is a collaborative project")

# substitute first 'R' with 'RR'
sub("R", "RR", Rstring)

## [1] "The RR Foundation"           "for Statistical Computing"
## [3] "RR is FREE software"        "RR is a collaborative project"
```

Replace all occurrences with `gsub()`

To replace not only the first pattern occurrence, but **all** of the occurrences we should use `gsub()` (think of it as *general* substitution)

Replacing with gsub()

```
# string
Rstring <- c("The R Foundation",
             "for Statistical Computing",
             "R is FREE software",
             "R is a collaborative project")

# substitute all 'R' with 'RR'
gsub("R", "RR", Rstring)

## [1] "The RR Foundation"           "for Statistical Computing"
## [3] "RR is FRREE software"       "RR is a collaborative project"
```

Replacing with `str_replace()`

```
# string
Rstring <- c("The R Foundation",
             "for Statistical Computing",
             "R is FREE software",
             "R is a collaborative project")

# substitute first 'R' with 'RR'
str_replace(Rstring, "R", "RR")

## [1] "The RR Foundation"           "for Statistical Computing"
## [3] "RR is FREE software"        "RR is a collaborative project"
```

Replacing with `str_replace_all()`

```
# string
Rstring <- c("The R Foundation",
             "for Statistical Computing",
             "R is FREE software",
             "R is a collaborative project")

# substitute first 'R' with 'RR'
str_replace_all(Rstring, "R", "RR")

## [1] "The RR Foundation"           "for Statistical Computing"
## [3] "RR is FRREE software"       "RR is a collaborative project"
```

Split a string

Split a string

Another common task is *splitting* a string based on a pattern. The idea is to split the elements of a character vector into substrings according to regex matches.

Split a string

Functions for extracting positions of matched pattern:

- ▶ `strsplit()`
- ▶ `str_split() ("stringr")`

Split a string

```
# a sentence
sentence <- c("The quick brown fox")

# split into words
strsplit(sentence, " ")

## [[1]]
## [1] "The"    "quick" "brown" "fox"
```

Split a string

```
# telephone numbers
tels <- c("510-548-2238", "707-231-2440", "650-752-1300")

# split each number into its portions
strsplit(tels, "-")

## [[1]]
## [1] "510" "548" "2238"
##
## [[2]]
## [1] "707" "231" "2440"
##
## [[3]]
## [1] "650" "752" "1300"
```

Split a string

```
# telephone numbers
tels <- c("510-548-2238", "707-231-2440", "650-752-1300")

# split each number into its portions
str_split(tels, "-")

## [[1]]
## [1] "510" "548" "2238"
##
## [[2]]
## [1] "707" "231" "2440"
##
## [[3]]
## [1] "650" "752" "1300"
```

Special Characters

Metacharacters

- ▶ The simplest form of regular expressions are those that match a single character
- ▶ Most characters, including all letters and digits, are regular expressions that match themselves
- ▶ For example, the pattern "1" matches the number 1
- ▶ The pattern "blu" matches the set of letters "blu".
- ▶ However, there are some special characters that don't match themselves
- ▶ These special characters are known as **metacharacters**.

Metacharacters

The metacharacters in *Extended Regular Expressions* are:

. \ | () [{ \$ * + ?

To use a metacharacter symbol as a literal character, we need to **escape** them

Metacharacters and how to escape them in R

Metacharacter	Literal meaning	Escape in R
.	the period or dot	\\.
\$	the dollar sign	\\\$
*	the asterisk or star	*
+	the plus sign	\\+
?	the question mark	\\?
	the vertical bar or pipe symbol	\\
\\	the backslash	\\\\
^	the caret	\\^
[the opening bracket	\\[
]	the closing bracket	\\]
{	the opening brace	\\{
}	the closing brace	\\}
(the opening parenthesis	\\(
)	the closing parenthesis	\\)

Character Classes

- ▶ A *character class* or *character set* is a list of characters enclosed by square brackets []
- ▶ Character sets are used to match **only one** of several characters.
- ▶ e.g. the regex character class [aA] matches any lower case letter a or any upper case letter A.
- ▶ character classes including the caret ^ at the beginning of the list indicates that the regular expression matches any character **NOT** in the list
- ▶ A dash symbol - (not at the beginning) is used to indicate ranges: e.g. [0-9]

Character Classes

```
# some string
transport <- c("car", "bike", "plane", "boat")

# look for 'e' or 'i'
grep(pattern = "[ei]", transport, value = TRUE)

## [1] "bike"  "plane"
```

Character Classes

Some (Regex) Character Classes

Anchor	Description
[aeiou]	match any one lower case vowel
[AEIOU]	match any one upper case vowel
[0123456789]	match any digit
[0-9]	match any digit (same as previous class)
[a-z]	match any lower case ASCII letter
[A-Z]	match any upper case ASCII letter
[a-zA-Z0-9]	match any of the above classes
[^aeiou]	match anything other than a lowercase vowel
[^0-9]	match anything other than a digit

POSIX Classes

Closely related to the regex character classes we have what is known as *POSIX character classes*. In R, POSIX character classes are represented with expressions inside double brackets `[[]]`.

POSIX Classes

POSIX Character Classes in R

Class	Description
<code>[:lower:]</code>	Lower-case letters
<code>[:upper:]</code>	Upper-case letters
<code>[:alpha:]</code>	Alphabetic characters (<code>[:lower:]</code> and <code>[:upper:]</code>)
<code>[:digit:]</code>	Digits: 0, 1, 2, 3, 4, 5, 6, 7, 8, 9
<code>[:alnum:]</code>	Alphanumeric characters (<code>[:alpha:]</code> and <code>[:digit:]</code>)
<code>[:blank:]</code>	Blank characters: space and tab
<code>[:cntrl:]</code>	Control characters
<code>[:punct:]</code>	Punctuation characters: ! " # % & ' () * + , - . / : ;
<code>[:space:]</code>	Space characters: tab, newline, vertical tab, form feed, carriage return, and space
<code>[:xdigit:]</code>	Hexadecimal digits: 0-9 A B C D E F a b c d e f
<code>[:print:]</code>	Printable characters (<code>[:alpha:]</code> , <code>[:punct:]</code> and space)
<code>[:graph:]</code>	Graphical characters (<code>[:alpha:]</code> and <code>[:punct:]</code>)

Special Sequences

Sequences define, no surprinsingly, sequences of characters which can match.

There are shorthand versions (or anchors) for commonly used sequences

Anchor Sequences in R

Anchor	Description
<code>\\d</code>	match a digit character
<code>\\D</code>	match a non-digit character
<code>\\s</code>	match a space character
<code>\\S</code>	match a non-space character
<code>\\w</code>	match a word character
<code>\\W</code>	match a non-word character
<code>\\b</code>	match a word boundary
<code>\\B</code>	match a non-(word boundary)
<code>\\h</code>	match a horizontal space
<code>\\H</code>	match a non-horizontal space
<code>\\v</code>	match a vertical space
<code>\\V</code>	match a non-vertical space

Quantifiers

- ▶ Another important set of regex elements are the so-called *quantifiers*.
- ▶ These are used when we want to match a **certain number** of characters that meet certain criteria.
- ▶ Quantifiers specify how many instances of a character, group, or character class must be present in the input for a match to be found.

Quantifiers

Quantifier	Description
?	The preceding item is optional and will be matched at most once
*	The preceding item will be matched zero or more times
+	The preceding item will be matched one or more times
{ <i>n</i> }	The preceding item is matched exactly <i>n</i> times
{ <i>n</i> , }	The preceding item is matched <i>n</i> or more times
{ <i>n</i> , <i>m</i> }	The preceding item is matched at least <i>n</i> times, but not more than <i>m</i> times

Positions

Character	Description
^	matches the start of the string
\$	matches the end of the string
.	matches any single character
	"OR" operator, matches patterns on either side of the

Regular Expressions in R

- ▶ There are two main aspects that we need to consider about regular expressions in R
- ▶ One has to do with *the functions* designed for regex pattern matching
- ▶ The other aspect has to do with *the way* regex patterns are expressed in R